

```
In [1]: # Import the necessary library for Data manipulation and visualization
import pandas as pd #for loading and data manipulation
import numpy as np #for mathematical operations
import matplotlib.pyplot as plt #for graphical representation of data
import seaborn as sns #for graphical representation of data
```

```
In [2]: df = pd.read_csv('C:/Users/Sanayak/Desktop/customer_purchase_data.csv')
print(df.head(10))#inspect the first 10 entries of our dataset
```

	Age	Gender	AnnualIncome	NumberOfPurchases	ProductCategory	\
0	40	1	66120.267939	8	0	
1	20	1	23579.773583	4	2	
2	27	1	127821.306432	11	2	
3	24	1	137798.623120	19	3	
4	31	1	99300.964220	19	1	
5	66	1	37758.117475	14	4	
6	39	1	126883.385286	16	3	
7	64	1	39707.359724	13	2	
8	43	0	102797.301269	20	1	
9	20	1	63854.921080	16	0	

	TimeSpentOnWebsite	LoyaltyProgram	DiscountsAvailed	PurchaseStatus
0	30.568601	0	5	1
1	38.240097	0	5	0
2	31.633212	1	0	1
3	46.167059	0	4	1
4	19.823592	0	0	1
5	17.827493	0	2	0
6	42.085384	1	4	1
7	17.190292	1	0	0
8	6.023475	0	3	0
9	38.572466	0	5	1

```
In [3]: print(df.duplicated().sum())#check for duplicate entries
```

112

```
In [4]: df.drop_duplicates(inplace=True)#drop the duplicated entries and update the dataset
print(df.duplicated().sum())
```

0

```
In [5]: print(df.info())#call the .info()function and inspect the data type
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 1388 entries, 0 to 1499
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Age                    1388 non-null   int64
1   Gender                 1388 non-null   int64
2   AnnualIncome           1388 non-null   float64
3   NumberOfPurchases     1388 non-null   int64
4   ProductCategory        1388 non-null   int64
5   TimeSpentOnWebsite     1388 non-null   float64
6   LoyaltyProgram         1388 non-null   int64
7   DiscountsAvailed       1388 non-null   int64
8   PurchaseStatus         1388 non-null   int64
dtypes: float64(2), int64(7)
memory usage: 108.4 KB
None
```

```
In [6]: df.describe().T#call the .describe()function to check for statistical information a
```

Out[6]:

	count	mean	std	min	25%	50%	75%	max
Age	1388.0	43.939481	15.487533	18.000000	30.750000	43.000000	54.000000	74.000000
Gender	1388.0	0.501441	0.500178	0.000000	0.000000	0.000000	1.000000	1.000000
AnnualIncome	1388.0	84699.045444	37541.136478	20001.512518	53766.895806	84699.045444	119999.999999	149999.999999
NumberOfPurchases	1388.0	10.548991	5.869383	0.000000	6.000000	10.000000	15.000000	20.000000
ProductCategory	1388.0	2.002882	1.422851	0.000000	1.000000	2.000000	3.000000	4.000000
TimeSpentOnWebsite	1388.0	30.747545	16.976852	1.037023	16.379635	30.747545	47.000000	60.000000
LoyaltyProgram	1388.0	0.333573	0.471659	0.000000	0.000000	0.000000	1.000000	1.000000
DiscountsAvailed	1388.0	2.609510	1.699984	0.000000	1.000000	2.000000	3.000000	4.000000
PurchaseStatus	1388.0	0.466859	0.499080	0.000000	0.000000	0.000000	1.000000	1.000000

The above statistical information shows that the Dataset has no outlier as the discrepancy between the mean and the 50thpercentile (median) can be ignored

```
In [7]: df[['AnnualIncome','TimeSpentOnWebsite']] = df[['AnnualIncome','TimeSpentOnWebsite']
print(df.head())#roundoff the Annualincome & Time spent on website column to 2decim
```

	Age	Gender	AnnualIncome	NumberOfPurchases	ProductCategory	\
0	40	1	66120.27	8	0	
1	20	1	23579.77	4	2	
2	27	1	127821.31	11	2	
3	24	1	137798.62	19	3	
4	31	1	99300.96	19	1	

	TimeSpentOnWebsite	LoyaltyProgram	DiscountsAvailed	PurchaseStatus
0	30.57	0	5	1
1	38.24	0	5	0
2	31.63	1	0	1
3	46.17	0	4	1
4	19.82	0	0	1

Data Visualization

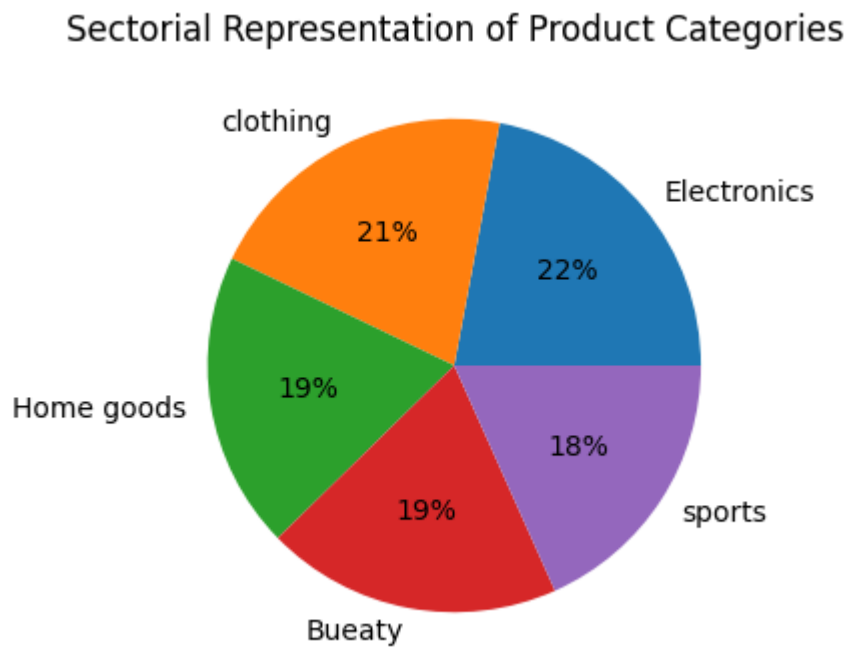
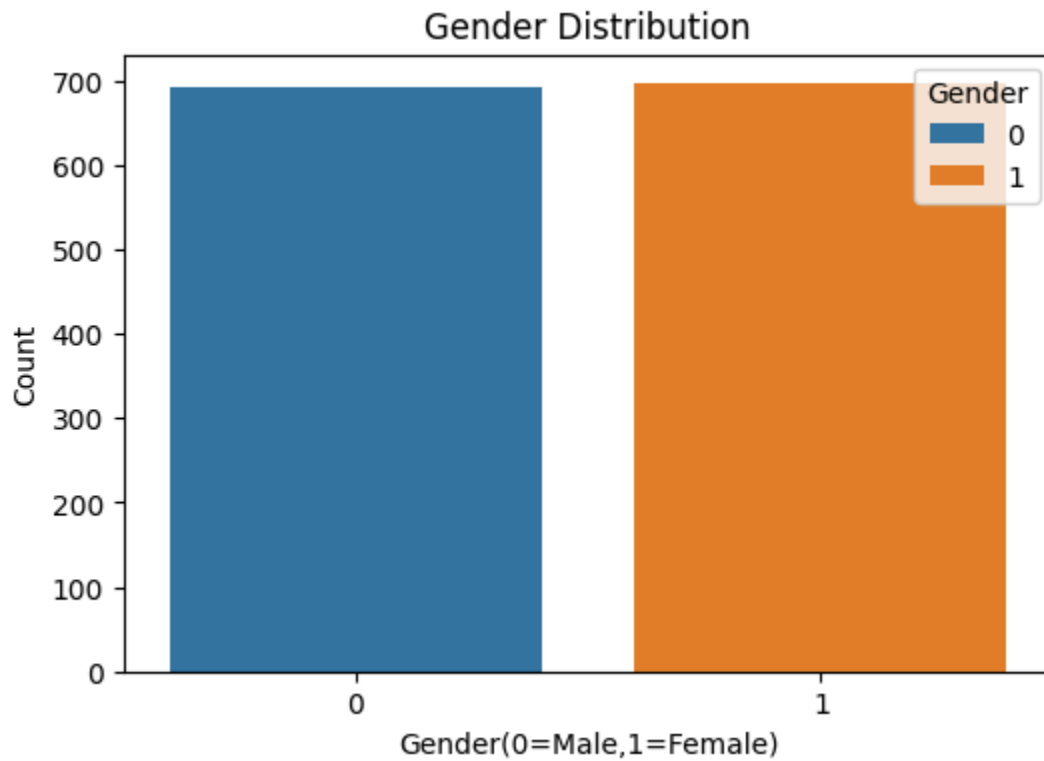
```
In [8]: # Gender Distribution
plt.figure(figsize=(6,4))
sns.countplot(x='Gender',data=df,hue='Gender')
plt.title('Gender Distribution')
plt.xlabel('Gender(0=Male,1=Female)')
plt.ylabel('Count')
plt.show()

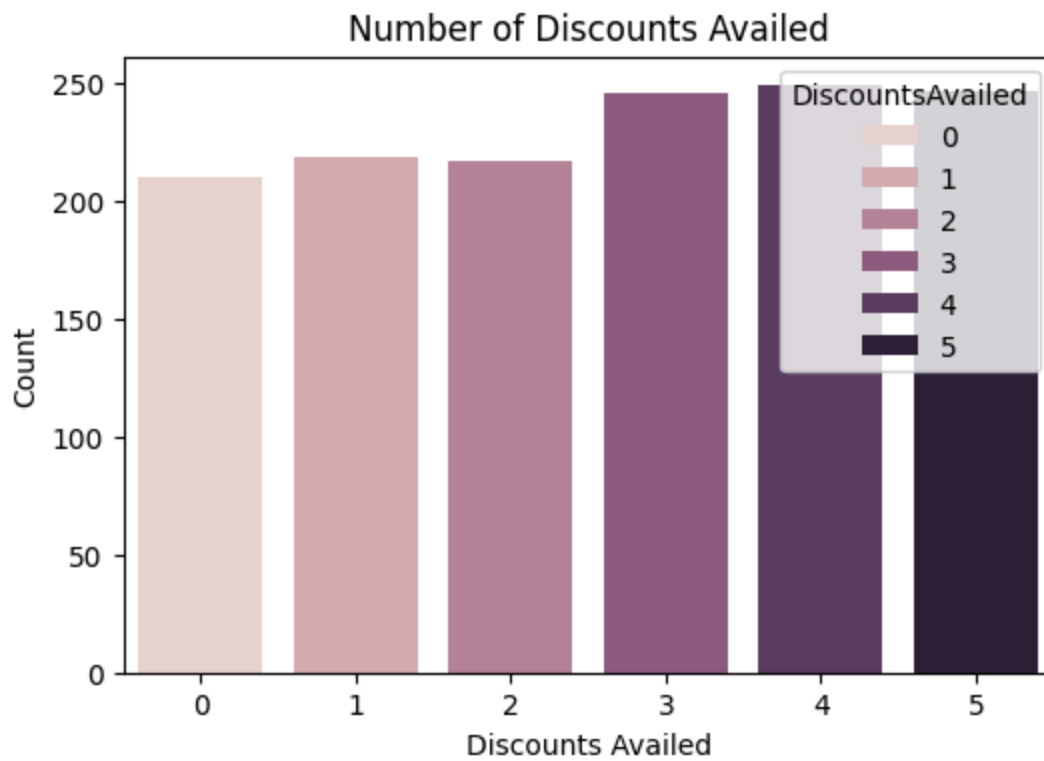
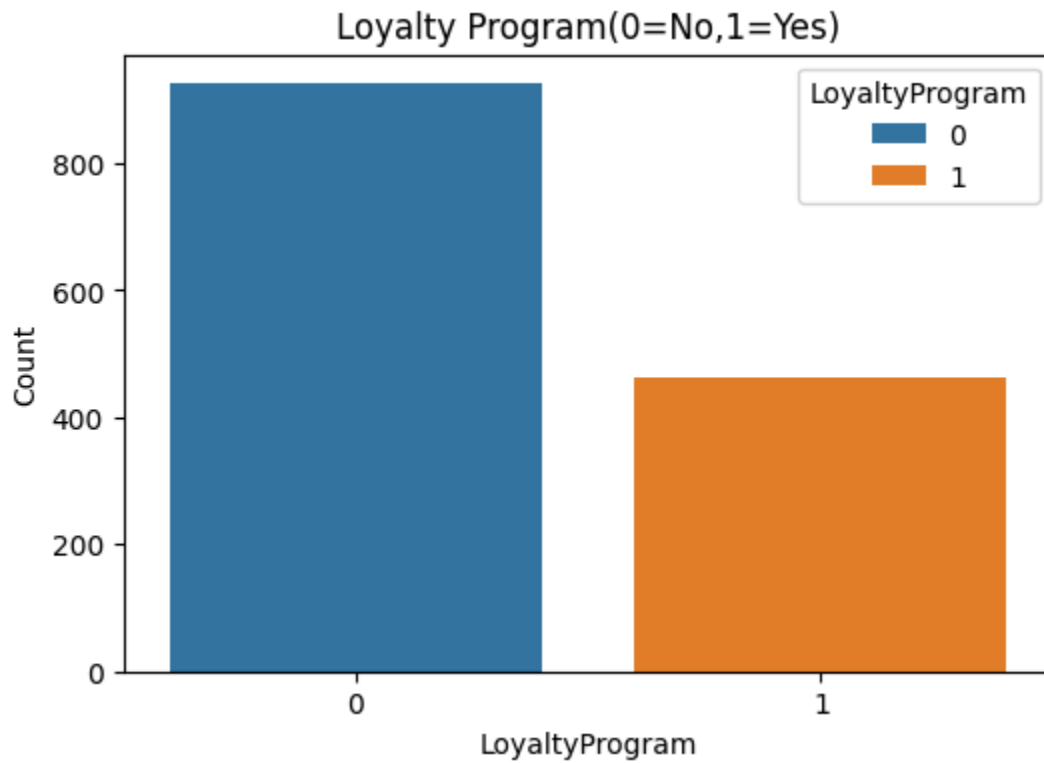
# Product Categories Distribution
product_cat_counts = df['ProductCategory'].value_counts()
labels = ['Electronics','clothing','Home goods','Bueaty','sports']
plt.figure(figsize=(6,4))
plt.pie(product_cat_counts,labels=labels,autopct="%1.f%%")
plt.title('Sectorial Representation of Product Categories')
plt.show()

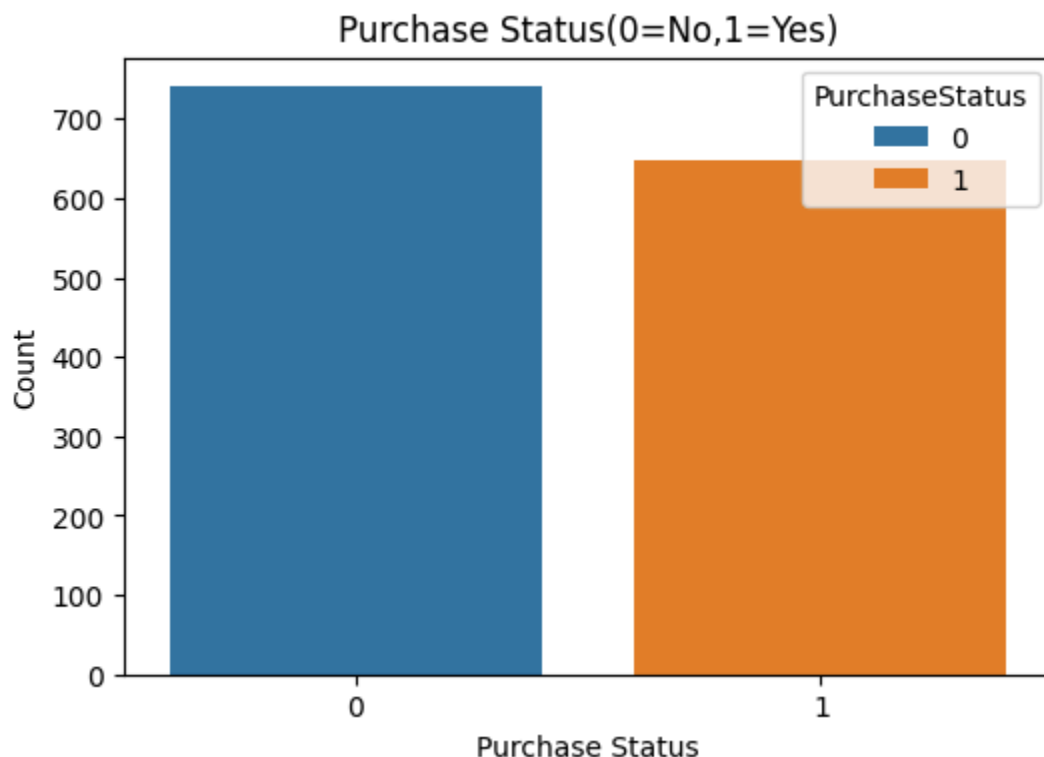
# Loyalty Program Distribution
plt.figure(figsize=(6,4))
sns.countplot(x='LoyaltyProgram',data=df,hue='LoyaltyProgram')
plt.title('Loyalty Program(0=No,1=Yes)')
plt.xlabel('LoyaltyProgram')
plt.ylabel('Count')
plt.show()

# Number of Discounts Availed Distribution
plt.figure(figsize=(6,4))
sns.countplot(x='DiscountsAvailed',data=df,hue='DiscountsAvailed')
plt.title('Number of Discounts Availed')
plt.xlabel('Discounts Availed')
plt.ylabel('Count')
plt.show()

# Purchase Status Distribution
plt.figure(figsize=(6,4))
sns.countplot(x='PurchaseStatus',data=df,hue='PurchaseStatus')
plt.title('Purchase Status(0=No,1=Yes)')
plt.xlabel('Purchase Status')
plt.ylabel('Count')
plt.show()
```





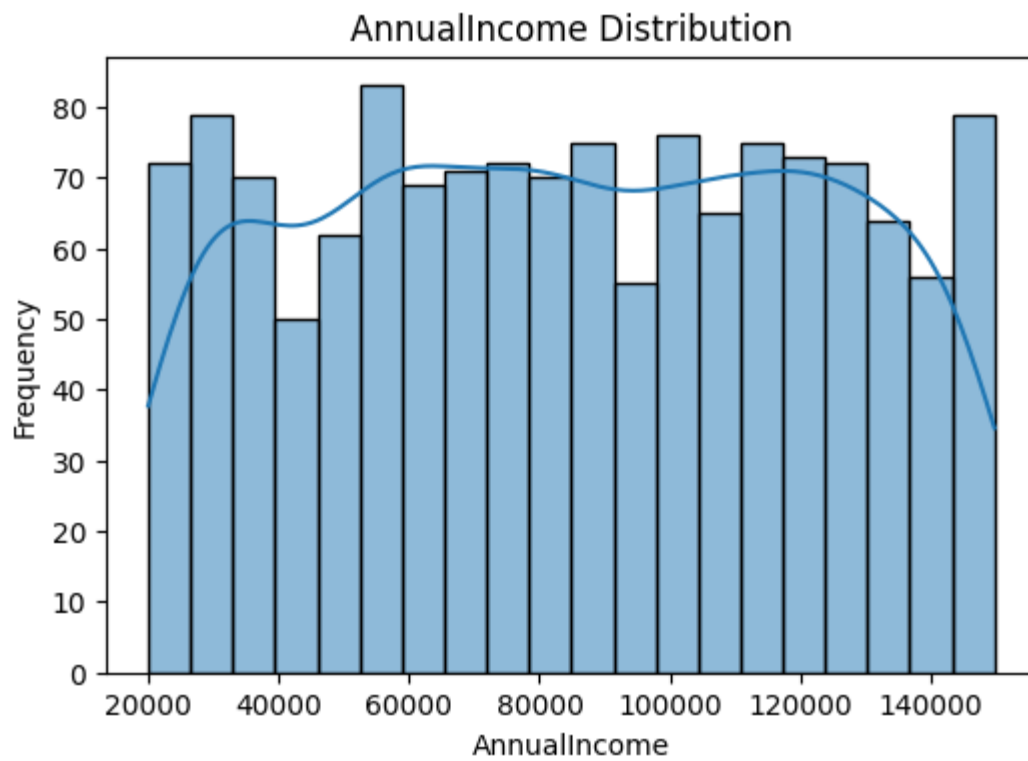



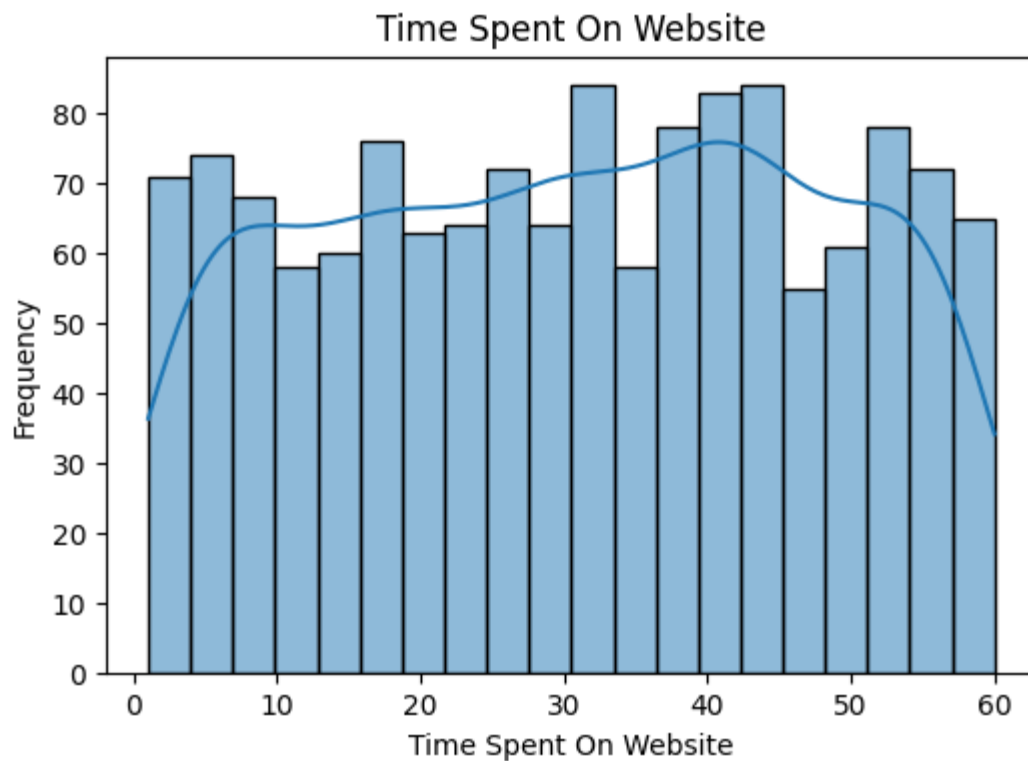

```
In [9]: # Age Distribution
plt.figure(figsize=(6,4))
sns.histplot(df['Age'],kde=True,bins=20)
plt.title('Age Distribution')
plt.xlabel('Age')
plt.ylabel('Frequency')
plt.show()

# Annual Income Distribution
plt.figure(figsize=(6,4))
sns.histplot(df['AnnualIncome'],kde=True,bins=20)
plt.title('AnnualIncome Distribution')
plt.xlabel('AnnualIncome')
plt.ylabel('Frequency')
plt.show()

# Number of Purchases Distribution
plt.figure(figsize=(6,4))
sns.histplot(df['NumberOfPurchases'],kde=True,bins=20)
plt.title('Number Of Purchases Distribution')
plt.xlabel('Number Of Purchases')
plt.ylabel('Frequency')
plt.show()

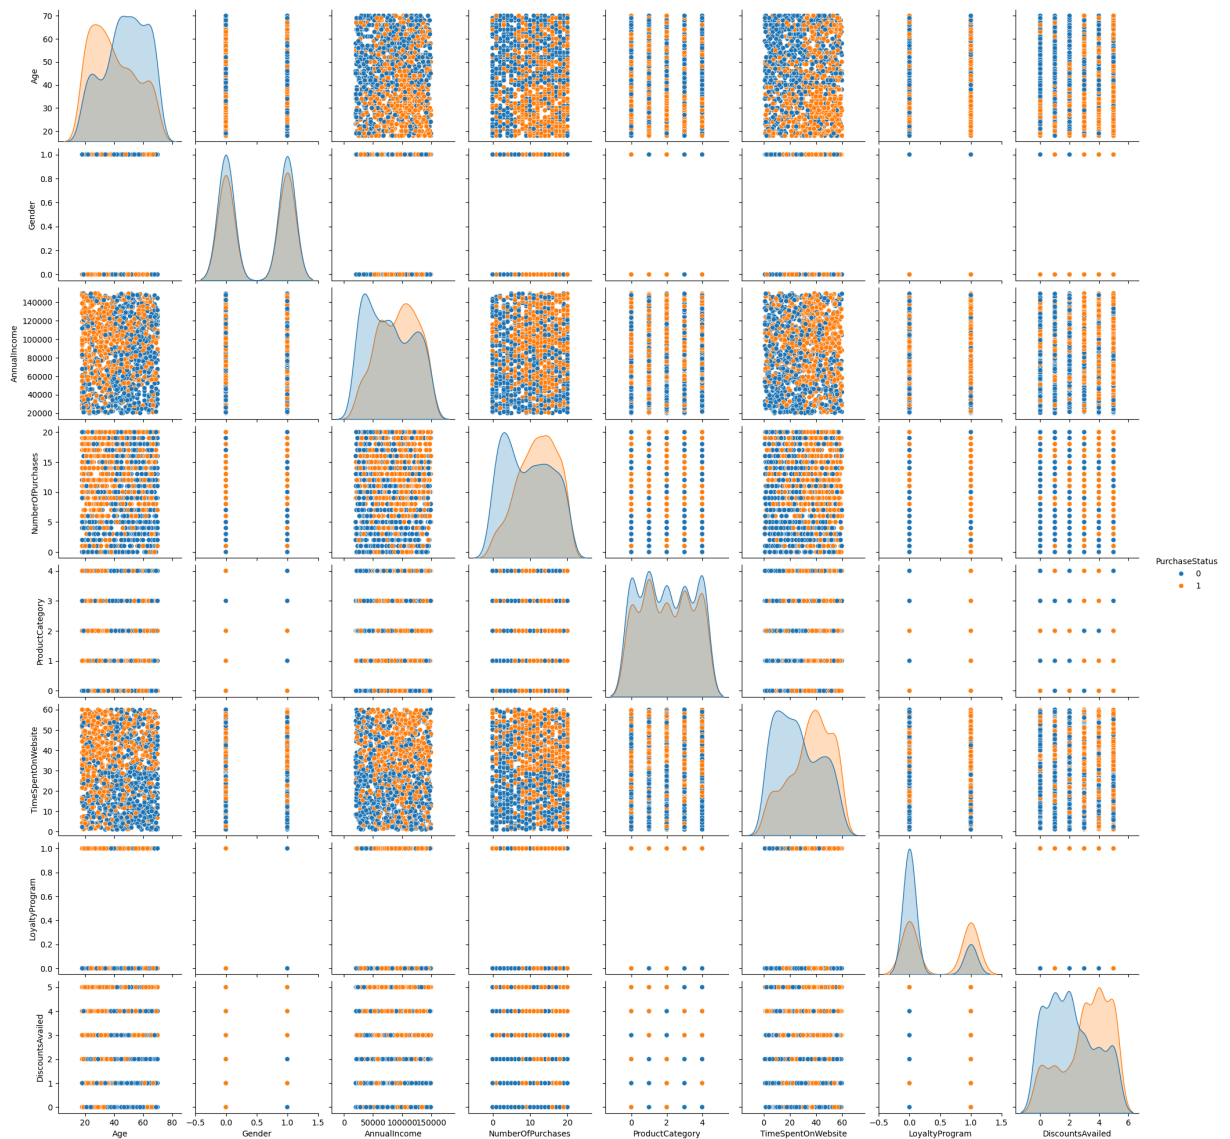
plt.figure(figsize=(6,4))
sns.histplot(df['TimeSpentOnWebsite'],kde=True,bins=20)
plt.title('Time Spent On Website')
plt.xlabel('Time Spent On Website')
plt.ylabel('Frequency')
plt.show()
```






```
In [10]: sns.pairplot(df, hue='PurchaseStatus')
```

```
Out[10]: <seaborn.axisgrid.PairGrid at 0x28f7cd38b00>
```

```
In [11]: # check for Multicollinearity
plt.figure(figsize=(6,5))
sns.heatmap(df.corr(),annot=True,cmap='coolwarm',fmt='.2f')
plt.title('Correlation Matrix')
plt.show()
```



The above Correlation matrix indicates that the features of the Dataset are not linearly dependent. This check is important as Multicollinearity among independent variables will result in less statistical inferences.

```
In [12]: # import the relevant module for our predictive model
from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.compose import ColumnTransformer
from sklearn.model_selection import train_test_split
from sklearn.ensemble import GradientBoostingClassifier
from sklearn.pipeline import Pipeline
from sklearn.metrics import classification_report, confusion_matrix, accuracy_score

In [13]: # Define feature columns and target column
X = df.drop('PurchaseStatus', axis=1)
y = df['PurchaseStatus']
# Identify numerical columns and categorical columns
cat_cols = ['Gender', 'ProductCategory', 'LoyaltyProgram', 'DiscountsAvailed', 'PurchaseStatus']
num_cols = ['Age', 'AnnualIncome', 'NumberOfPurchases', 'TimeSpentOnWebsite']

# create a ColumnTransformer for preprocessing
```

```

preprocessor = ColumnTransformer(transformers=[('num',StandardScaler(),num_cols),
                                              ('cat',OneHotEncoder(),cat_cols[:-1])]
                                # create a pipeline that first preprocess the data then applies the model
                                pipeline = Pipeline([('preprocessor',preprocessor),
                                                    ('classifier',GradientBoostingClassifier(random_state=42))])

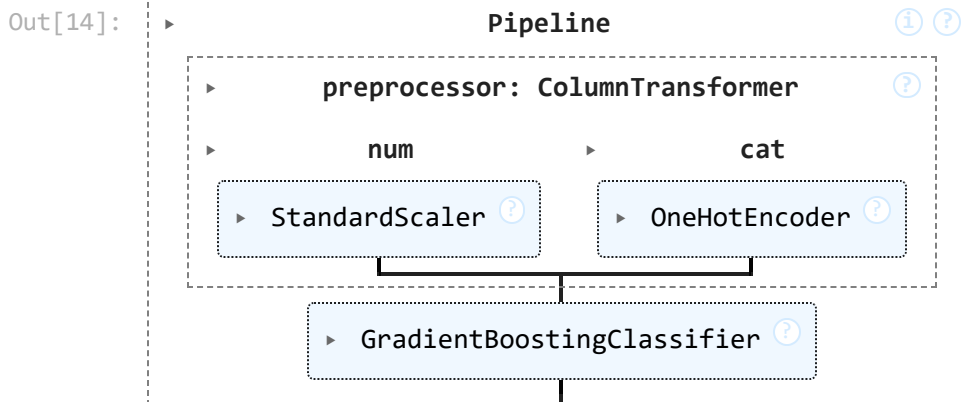
```

```

In [14]: # Split the data into training and testing stes with train size of 75% of our data
X_train,X_test,y_train,y_test = train_test_split(X,y,
                                                  train_size = .75,
                                                  stratify = y,
                                                  random_state = 42)

# Train the model
pipeline.fit(X_train,y_train)

```



```

In [15]: # Make predictions
y_pred = pipeline.predict(X_test)
# Evaluate model performance on test set
accuracy = accuracy_score(y_test,y_pred)
class_report = classification_report(y_test,y_pred)
cm = confusion_matrix(y_test,y_pred)

```

```

In [16]: # Output the results
print('Model Test Accuracy:',accuracy)
print('Classification Report:\n',class_report)


# Visualize the confusion matrix
plt.figure(figsize=(6,4))
sns.heatmap(cm, annot = True,
            fmt='d',cmap='Blues',
            xticklabels=(['Predicted Negative','Predicted Positive']),
            yticklabels=(['Atual Positive','Actual Negative']))
plt.title('Confusion Matrix')
plt.xlabel('Predicted Label')
plt.ylabel('True Label')
plt.show()

```

Model Test Accuracy: 0.9106628242074928

Classification Report:

	precision	recall	f1-score	support
0	0.91	0.93	0.92	185
1	0.92	0.89	0.90	162
accuracy			0.91	347
macro avg	0.91	0.91	0.91	347
weighted avg	0.91	0.91	0.91	347



PREDICTING CUSTOMER'S BEHAVIOUR

Description:

This dataset contains information on customer purchase behavior across various attributes, aiming to help understand the factors influencing purchase decisions. The dataset includes purchasing habits and other relevant features.

Features:

Age: Customer's age

Gender: Customer's gender (0: Male, 1: Female)

Annual Income: Annual income of the customer in dollars

Number of Purchases: Total number of purchases made by the customer

Product Category: Category of the purchased product (0: Electronics, 1: Clothing, 2: Home Goods, 3: Beauty, 4: Sports)

Time Spent on Website: Time spent by the customer on the website in minutes

Loyalty Program: Whether the customer is a member of the loyalty program (0: No, 1: Yes)

Discounts Availed: Number of discounts availed by the customer (range: 0-5)

Purchase Status (Target Variable): Likelihood of the customer making a purchase (0: No, 1: Yes)

ANALYZING THE MODEL PERFORMANCE

Model Test Accuracy: The model correctly predicted 91% of the test samples. This is a relatively high accuracy, indicating that the model is performing well overall.

Precision: Measures the proportion of correct predictions among those predicted as a class. For class 0: 91% of instances predicted as 0 were actually 0. For class 1: 92% of instances predicted as 1 were actually 1.

Recall: Measures the proportion of correctly predicted instances among the actual instances of a class. For class 0: 93% of actual 0 instances were correctly predicted. For class 1: 89% of actual 1 instance were correctly predicted.

F1-score: The harmonic mean of precision and recall, balancing both metrics.

Support: The number of instances for each class.

Overall Conclusion: The model exhibits strong performance, achieving an overall accuracy of 91%. Both classes demonstrate high precision and recall, indicating that the model is able to accurately identify instances of both classes. The macro average and weighted average of precision, recall, and F1-score further reinforce the model's balanced performance across both classes.

Key Takeaways: The model is well-suited for this classification task.

METHODOLOGY

Tools & Libraries used for Analysis:

Pandas: For Data manipulation and

Numpy: |For Numerical Operations

Matplotlib & Seaborn: For Graphical representation of the Dataset.

Scikit-learn: The GradientBoostingClassifier class of Scikit-learn which is an Ensemble Machine learning technic for model classification is employed.

Reference material: <https://www.kaggle.com/datasets/rabieelkharoua/predict-customer-purchase-behavior-dataset>